

Tips:

- Communicate with Kostas
- Use figures!
- Ensure structure is flowing
- Use sources for everything
- Modularise in order to convert to paper

To Do:

- Modularise in order to convert to paper
- Describe appendix
- Act on Interim Review Feedback
- Act on planning and timescale comments from interim report
- Act on contents of summary of completed work from interim report
- Change the Rationale to fit Final report
- Write a literature review or combine with Rationale and add to it?
- Write Contents and Knowledge
- Write critical analysis and discussion
- Write professional issues (get background reading and citations)

Final Year Project Report

Full Unit - Final Report

Cooperative Strategies in Multi-Agent Systems

James King

A report submitted in part fulfilment of the degree of

BSc (Hons) in Computer Science

Supervisor: Kostas Stathis



Department of Computer Science
Royal Holloway, University of London

February 5, 2019

Declaration

This report has been prepared on the basis of my own work. Where other published and unpublished source materials have been used, these have been acknowledged.

Word Count:

Student Name: James King

Date of Submission: February 5, 2019

Signature:

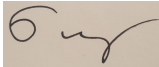
A handwritten signature in black ink on a light brown rectangular background. The signature is cursive and appears to read 'James King'.

Table of Contents

Abstract	4
1 Introduction	5
1.1 Motivation	5
1.2 Aims and Objectives	5
1.3 Contribution	5
1.4 Structure	5
2 Background	6
2.1 Introduction	6
2.2 Technique 1	6
2.3 Technique 2	6
2.4 Technique n	6
2.5 Summary	6
3 Framework	7
3.1 Theoretical	7
3.2 Implementation	7
3.3 Technical Preparation	7
3.4 Design	8
3.5 Development	8
3.6 Testing	8
3.7 Software Engineering	8
3.8 Documentation	8
3.9 Experiment Evaluation	8
4 Critical Analysis and Discussion	10
5 Conclusions	11

Bibliography 12

6 Professional Issues 13

7 Appendix 14

Abstract

Chapter 1: **Introduction**

1.1 Motivation

Problem statement and associated objectives.
Give a rationale.

1.2 Aims and Objectives

1.3 Contribution

1.4 Structure

Chapter 2: **Background**

2.1 **Introduction**

Include background reading
Critical lit review.
Comprehensive knowledge.

2.2 **Literature Review**

2.3 **Technique 1**

Description, limitations, our technique
Examples:

- Evolcoop Axelrod, Axelrod.py
- Arithmetics of mutual help
- Phekpls, game theoretic analysis
- 5 rules coop
- spatial and graphs
- Evol indirect image, Nowak and Sigmund
- multilevel nowak
- Genetic drift as an issue
- gossip alternative
- rational cooperation in finite iterated prisoners
- prisoners and their dilemma
- evol dir indir
- evol of coop through indir rec
- kin hamilton
- image vs standing
- notions of reputation in mas

2.4 Technique 2

Description, limitations, our technique

2.5 Technique n

Description, limitations, our technique

2.6 Summary

Summarise limitations and the techniques we are using to combat this

Chapter 3: **Framework**

Excellent understanding and insight. Conceptual framework underpins study. Comprehensive expert account of topic. Well thought through Software Engineering content.

3.1 Theoretical

Relate to background.

3.2 Implementation

Relate to agents systems.

3.2.1 Agents Service

3.2.2 Environment

3.2.3 Web Application and Interface

3.2.4 Conclusion

3.3 Technical Preparation

Learning flask etc.

3.4 Design

3.4.1 Agents Service

3.4.2 Environment

3.4.3 Web Application and Interface

3.5 Development

3.5.1 Agents Service

3.5.2 Environment

3.5.3 Web Application and Interface

3.6 Testing

3.6.1 Agents Service Testing

3.6.2 Environment Testing

3.6.3 Web Application Testing

3.7 Software Engineering

3.8 Documentation

Software documentation, user guide.

3.9 Experiment Evaluation

3.9.1 Experiment 1

Variable setting, controls etc.
Evaluation

3.9.2 Experiment 2

3.9.3 Experiment n

3.9.4 Discussion

Evaluation and discussion on the experiments

Successful?

Opportunity for extra work?

Focus on experiments

Chapter 4: **Critical Analysis and Discussion**

Project achievements

Reflection on project process (difficulties, success, failure)

Successful?

Future enhancements

Focus on project

Compare to deliverables at beginning.

Chapter 5: **Conclusions**

Opportunity for extra work?

Bibliography

1

¹A lot of my background theory work has been completed in my earlier reports so many of the references appear in those reports (see appendix)

Chapter 6: **Professional Issues**

Replacement of people in jobs?

AI issues and risks?

Relate to project

Chapter 7: **Appendix**

Describe the contents of my appendix.

Report on the evolution of cooperation in relation to game-theory and different game-theoretic mechanisms to aid it

JAMES KING

Supervisor: Kostas Stathis

October 2018

Abstract

Since the early days of Darwinian evolutionary theory the phrase "Survival of the fittest" has become synonymous with much thinking on evolutionary dynamics. This idea has come along way since then but is still seen by many to promote selfish attitudes. However, we can see throughout nature that cooperation between biological agents is prevalent. So how did cooperation evolve and why has it flourished? This question fundamentally challenges what seemed concrete ideas about evolutionary dynamics, and has been approached from many angles. Game-theory is a branch of mathematics that has spawned many mathematical models of evolutionary dynamics in an attempt to solve the problem, some have been formulated programmatically to analyse the results. Study of the evolution of cooperation also has wider impacts in the world of computer science - namely on agent-based systems. A large component of agent-based system design is how interactions between agents works and how cooperation can be garnered within societies of agents. In this report, I shall explore past game-theoretic approaches to this problem. This exploration has helped me gather a deeper understanding of evolutionary dynamics, the reasoning behind these mathematical approaches to the problem of the evolution of cooperation and their application to the area of intelligent agents and multi-agent systems.

I. INTRODUCTION

Early Darwinians often focused on the struggle for survival as a means to explain the phenomena of evolution. This focus falls short on explaining the phenomena of altruistic behaviour, highlighted early on by Kropotkin [10]. The phenomena presented a problem for evolutionary biologists, how could it have evolved from an inherently selfish world?

Altruistic acts are actions where an individual puts others above themselves, these acts benefits recipients and are at a cost in some way to the actor. Examples of these acts and behaviours have been well documented and pervade both the natural world and human society.

In the seminal paper on the evolution of cooperation [1] Axelrod and Hamilton identify two theories proposed to solve the problem: kinship theory and reciprocation theory, focusing on the latter - particularly the Iterated Prisoner's Dilemma.

Here I will explore these mechanisms that attempt to explain altruistic behaviour and I will endeavour to relate them to multi-agent systems or even point out how they are inapplicable to this domain if appropriate.

II. CONTENT AND KNOWLEDGE

i. Kinship Theory

Kinship theory is an umbrella term for a number of models that attempt to solve the problem of the evolution of cooperation. The tie between these theories is that the individuals who choose to cooperate with each other are in some way ‘kin’. The definition of what ‘kin’ is varies depending on the theory and the purpose for that theory.

Richard Dawkins popularized the idea of the ‘selfish gene’ [2]. This idea argues that as genes are the actual replicators, actors are hardwired to propagate the gene. This propagation does not only involve reproduction but also acts of cooperation to support and maintain others who share the gene. This is evident in many areas of nature especially in family groups, shown even recently in the BBC series *Dynasties*.

W.D. Hamilton supports a similar view [4]: that individuals with shared genes will work together to preserve those genes. He argues that to do this individuals don’t work to add to their own fitness, but work to improve their ‘inclusive fitness’, which includes the fitness of other related individuals. His model converts the extent of the relatedness of an individual to a quantitative value using Wright’s Coefficient of Relationship, which is shown in figure 1 on page 2. The model then uses this coefficient of relatedness to calculate inclusive fitness.

Hamilton finds that for that the value for the coefficient of relatedness must be greater than the cost-to-benefit ratio of the altruistic act for cooperation to evolve: $r > c/b$.

Axelrod and Hamilton [1] highlight that although the theory works well to explain cooperation between related individual it falls short in explaining the evolution of cooperation between unrelated individuals. This issue also makes it difficult to apply kinship theory as an aid to the evolution of cooperation in multi-agent systems. Kinship theory requires a reproductive drive and metrics for relatedness in order to motivate altruistic acts, both of which don’t relate closely to multi-agent




Figure 1: Wright’s coefficient of relatedness by Citynoise - Own work, CC BY-SA 4.0, <https://commons.wikimedia.org/w/index.php?curid=37723128>

systems though in some applications could be present.

ii. The Iterated Prisoner’s Dilemma

The Iterated Prisoner’s Dilemma is one such example of a game-theoretic model that uses direct reciprocity to attempt to solve the problem that the evolution of cooperation puts forward. Direct reciprocity is the idea: “If I cooperate now, you may cooperate later” [7]. The Dilemma is, in fact, a game played between two individuals. The game is made up of a number of repeated rounds, in each round the two players can choose between cooperating with or defecting against each other.

A payoff matrix is provided such as the one in figure 1 on page 3. These matrices provide a temptation to defect, but a reward if both players cooperate over multiple rounds. In a single round game it is mathematically best to defect, but in repeated rounds agents are encouraged to cooperate by a mutual gain in score. Nowak [7] reports that the cost-to-benefit ratio of the altruistic act must be less than the probability of another encounter for cooperation to evolve: $w > c/b$.

The dilemma has been extended into tourna-

ments such as the one in Axelrod and Hamilton [1]. You can even run tournaments directly in Python using the Axelrod-Python library [3]. Tournaments can come in a variety of styles but the most popular is the round-robin tournament where every player plays a match of the iterated prisoner's dilemma against every other player. Players accumulate points throughout these games, based on the payoff matrix.

This has been even further extended to include a genetic algorithm to simulate evolution by reproducing players into the next generation. Multiple different algorithms have been used, one common one is the Moran Process (A stochastic process used to model evolution in a finite, unstructured population) available in the Axelrod-Python library.

Player A	Player B	
	Cooperation	Defection
Cooperation	A=3 B=3	A=0 B=5
Defection	A=5 B=0	A=1 B=1

Table 1: The payoff matrix in a typical iterated prisoner's dilemma game (such as Axelrod and Hamilton's). $A=x$, $B=y$ where x denotes the payoff for A and y the payoff for B.

iii. Strategies

Players in the iterated prisoner's dilemma have to be able to decide whether to defect or cooperate in any given round against any given opponent. The player has access to the history of interaction between the two individuals to base their decision on - the simplest of strategies, however, don't use the history, one example being a pure defector.

Some more interesting strategies include the world famous tit-for-tat, grudger and Pavlov (win-stay, lose-shift [6]). The classical form of tit-for-tat begins by cooperating in the first round and then copying the other player's last move, this strategy is so effective because it

gives the other player a chance to cooperate (so working well with cooperators) but still punishes a defecting player preventing them from taking advantage. There are many other versions of tit-for-tat, a good example being forgiving tit-for-tat which only defects if the other player defects for two turns in a row.

Grudger works by beginning to cooperate but if the other player defects the individual switches and defects for the rest of the time, this is not so effective as tit-for-tat because it never forgives the other player, preventing further cooperation.

Pavlov works by continuing to act the way it has in the previous interaction if the action was successful if it is not successful the strategy switches to the other action. Nowak and Sigmund [6] claim that Pavlov outperforms tit-for-tat as tit-for-tat is unsuccessful in non-deterministic environments (such as one that has a random chance that an action will not be what a player chooses). The real world is non-deterministic and so tit-for-tat doesn't generalise into it so well. Pavlov considers a success in a round to be: you both cooperate or you defect and the other player cooperated, and considers a failure to be: you both defect or you cooperate and the other player defects.

iv. Other Reciprocation Theories

Using reciprocity to aid in the evolution of cooperation has not been limited to the iterated prisoner's dilemma. Martin A. Nowak presents 3 other mechanisms that use reciprocation [7]: indirect reciprocity, network reciprocity (also known as spatial tournaments) and group selection. A graphical description of the 5 rules for the evolution of cooperation is provided in figure 2.

Indirect Reciprocity is a set of mechanisms that leverage the group mechanic of reputation. Each member of the set of mechanisms uses different metrics for reputation and the spread of information that influences reputation. Nowak and Sigmund's version uses the idea of image scoring [9]. An image score represents how well an agent is thought of, the

higher the better.

Which mechanism is more effective is the subject of extensive debate, in two separate papers both Leimar and Hammerstein [5], and Roberts [11] concluded that the standing strategy is more effective. In the standing strategy agents do not have an image score but have a standing which can be good or bad, everyone starts off with a good standing but if they defect against a player with good standing they are considered to have a bad standing.

Nowak [7] reports that the probability of knowing someone's reputation must be greater than the cost-to-benefit ratio of the altruistic act for cooperation to evolve: $q > c/b$. How information is spread throughout a population playing a game of indirect reciprocity thus becomes important.

Though not all agents can realistically be on-lookers, another mechanism for the spread of information is through gossip [12]. Gossip can act as a mechanism for knowing the reputation of others, but can also muddy the water as to whether agents know the "correct" reputation of the recipient, or have they received "incorrect" information.

Network reciprocity (also known as spatial tournaments) uses direct reciprocity but not in a round-robin tournament. Spatial tournaments use a graph where nodes are agents and the edges represent channels of connection between agents [13], direct reciprocity games are then played between agents with connections. Nowak and May [8] found that groups of co-operators can work together to support each other and fend off invasion by defectors - depending on the structure of the graph.

Direct Reciprocity is also used in group selection. In this model the population is subdivided into smaller groups, in these smaller groups individuals interact with each other and reproduction occurs from the fitness they gain in these interactions.

Though offspring are reproduced into the same group, reproduction occurs on two levels: group sizes fluctuate due to reproduction, if a group gets to a certain size it may split in two and when a group divides another is




Figure 2: A graphical representation of the five rules of cooperation from Nowak [7]

eliminated to keep the number of groups fixed. Groups with faster reproduction do better than those with slower reproduction. Traulsen and Nowak found this favours the evolution of cooperation under certain conditions [14].

III. DISCUSSION AND CONCLUSION

All these mechanisms to aid in the evolution of cooperation have extremely interesting properties and consequences in both biological and intelligent agents. Network Reciprocity can be used to represent a physical telecommunications network across which agents communicate, group selection is useful to model the spread of successful agent populations displacing others and kinship theory is useful to model interactions between different family groups in nature just to name a few applications.

With the internet being such a vast medium across which agents may interact it is very likely that agents may not have many repeated

meetings, so direct reciprocity is not always applicable. However, agents may often repeat interactions, so indirect reciprocity is not always applicable.

Due to the incredible interconnectedness of the internet, the graphs used by network reciprocity would have to be highly connected and this makes it less useful for modelling agent systems using the internet. Due to the lack of general reproductive drive and inherent relatedness between agents across the internet Kinship theory is unlikely to apply well to all but exceptional agent systems.

Thus as my project will be aiming to simulate a multi-agent system with agents decentralised across the internet a combination of indirect and direct reciprocity will likely be the most germane aid to the evolution of cooperation for my project.

REFERENCES

- [1] Robert Axelrod and William D. Hamilton. The evolution of cooperation. *Science*, 211:1390–1396, 1981.
- [2] Richard Dawkins. *The Selfish Gene*. Oxford University Press, 2016.
- [3] Vince Knight; Owen Campbell; Marc; eric-s-s; VSN Reddy Janga; James Campbell; Karol M. Langner; T.J. Gaffney; Sourav Singh; Nikoleta; Julie Rymer; Thomas Campbell; Jason Young; MHakem; Geraint Palmer; Kristian Glass; edouardArgenson; Daniel Mancia; Martin Jones; Cameron Davidson-Pilon; alajara; Ranjini Das; Marios Zoulias; Aaron Kratz; Timothy Standen; Paul Slavin; Adam Pohl; Jochen MÄijller; Georgios Koutsovoulos; Areeb Ahmed. Axelrod: 4.3.0. <http://dx.doi.org/10.5281/zenodo.1405868>, September 2018.
- [4] W. D. Hamilton. The genetical evolution of social behaviour. *Journal of Theoretical Biology*, 7:1–16, July 1964.
- [5] Olof Leimar and Peter Hammerstein. Evolution of cooperation through indirect reciprocity. *Proceedings of The Royal Society*, 268:745–753, May 2001.
- [6] Martin Nowak and Karl Sigmund. A strategy of win-stay, lose-shift that outperforms tit-for-tat in the prisoner’s dilemma game. *Nature*, 364:56–58, July 1993.
- [7] Martin A. Nowak. Five rules for the evolution of cooperation. *Science*, 314, 2006.
- [8] Martin A. Nowak and Robert M. May. Evolutionary games and spatial chaos. *Nature*, 359:826–829, 1992.
- [9] Martin A. Nowak and Karl Sigmund. Evolution of indirect reciprocity by image scoring. *Nature*, 393:573–577, 1998.
- [10] Martin A. Nowak, Karl Sigmund, and Robert M. May. The arithmetics of mutual help. *Scientific American*, 272:76–86, 1995.
- [11] Gilbert Roberts. Evolution of direct and indirect reciprocity. *Proceedings of The Royal Society*, 275:173–179, September 2008.
- [12] Ralf D. Sommerfeld, Hans-Jürgen Krambeck, Dirk Semmann, and Manfred Milinski. Gossip as an alternative for direct observation in games of indirect reciprocity. *Proceedings of the National Academy of Sciences of the United States of America*, 104:17435–17440, 2007.
- [13] György Szabó and Gábor Fáth. Evolutionary games on graphs. *Physics Reports*, 446:96–216, 2007. Section: The structure of social graphs.
- [14] Arne Traulsen and Martin A. Nowak. Evolution of cooperation by multilevel selection. *Proceedings of the National Academy of Sciences of the United States of America*, 103:10952–10955, 2006.

Report on indirect reciprocity, strategies for agents and the development of a concrete model to implement

JAMES KING

Supervisor: Kostas Stathis

October 2018

Abstract

Indirect reciprocity is a mechanism that uses reciprocation theory to aid in the evolution of cooperation. Cooperation being an action taken by an individual that benefits another individual but at a cost to itself. Indirect reciprocity is a promising motivator for cooperation in societies with agents of higher intelligence levels and of greater sizes, such as human societies or even multi-agent systems. I plan to implement the mechanism programmatically, but there are many models and many possible additions to these models. In this report I will explore past approaches to indirect reciprocity, comparing and contrasting the variations proposed. The outcome of this exploration of approaches will be a concrete model to implement in this project.

I. INTRODUCTION

The evolution and preservation of cooperation has been a puzzle for evolutionary theorists for a long time. Many different approaches have been taken to give an explanation to cooperative phenomena, especially from the field of game theory. Often these approaches have come from the idea of reciprocity, where agents can grow mutually beneficial relationships through repeated interactions. The most popular mechanism of this being direct reciprocity where interactions are repeated between the same two individuals, and thus they may reciprocate directly with each other.

There is another game-theoretic mechanism known as indirect reciprocity, which works on the idea that nice agents will help those who help each other - "if I scratch your back, others will hopefully scratch mine" displayed graphically in figure 1. This means learning information about another player does not require direct interaction. A number of models have




Figure 1: The main idea of indirect reciprocity from Nowak [4]

been proposed to run indirect reciprocity. It is these models I shall be describing and reviewing, before using them to formulate my own to implement in my project.

II. REVIEW OF PAST WORK

i. Nowak, Sigmund and Image Scoring

Nowak and Sigmund's model has a population structure of just one population of players per generation, the first generation of players are chosen at the start of an indirect reciprocity tournament [5]. A population then reproduces into the next generation of players (more on reproduction further down). From the population of a generation a set of pairs of players are chosen at random to interact. A pair is made up of a donor and recipient. The donor can choose to cooperate at a cost c , from this the recipient receives b points ($b > c$), or defect at no cost to the donor's fitness but the recipient also receives no points.

The model utilizes the idea of image scoring. An image score is comparable to the idea of a player's reputation. A player's image score increases to those who see them cooperate and decreases to those who see them defect. The actual image score is an integer, bounded between -5 and 5. One key part of the model is where these image scores are held.

In the first formulation of this all image scores are public, but later on, it is suggested that this model is unlikely to hold true in real life as individual's views on other individuals' reputations tend to differ. As an alternative, the idea of onlookers is suggested, where a set of onlookers are chosen from the generation for each interaction. In both models, image scores are increased by one for a cooperation as a donor or decreased by one for a defection as a donor, though in the second model only the onlookers, recipient and donor can add this to their image score for the donor. The use of onlookers does make cooperation harder to establish, as it is dependent on whether a player knows an accurate image score of the recipient. The points accrued through interactions gives the fitness score of each player. The amount a player reproduces into the next generation is dependent on the fitness score. Mutation is also an option in this model, this phenomenon

is when a player reproduces at random their offspring can be a different strategy. Mutation can lead to loops where defectors are invaded by discriminators and then discriminators are undermined by cooperators - allowing defectors to take hold again.

The two simplest strategies are pure defection and pure cooperation. A key strategy for the evolution of cooperation is the discriminator strategy as illustrated in figure 2. In fact, there is a baseline number of discriminators required for cooperation to evolve. This strategy is given a number k , if the image score of a player is greater than or equal to k a discriminator will cooperate with them. More interesting strategies take into account their own image score or theirs and the recipient's image score, these can aid in the evolution of cooperation too. In the absence of information on other players, these strategies can believe other players to have an image score of k with a certain probability.

The model is known to be dependent on the probability of knowing the image of another player. Cooperation can only be stable if $q > c/b$, q being the probability of knowing the image, c being the cost and b being the benefit.

Leimar and Hammerstein recognised the problems of genetic drift in this model [2]. Genetic drift being "a situation in which the frequency of a particular gene in a small population of living things changes without a known cause" [1]. Adopting an island population model to restrict genetic drift. Rather than one single population group, in the island population model, the group is divided up into g amount of groups each with population n . This model also uses a different reproductive strategy where relative reproductive success locally is calculated as a normalisation of the sum of the fitness of all strategies within the same group. For global expected reproductive success sum each strategy and normalize over the whole population. A new individual is then locally derived with probability p or globally derived with probability $1-p$. The strategy for the individual produced is randomly generated with a distribution corresponding to the




Figure 2: How an image scoring discriminator strategy acts

local or global reproductive success.

ii. Standing Strategy

The standing strategy was suggested by both Leimar and Hammerstein [2] and Roberts [7] to be superior to image scoring. Their models used the island population structure given above. Leimar and Hammerstein replaced image scoring with the standing strategy whereas Roberts mixed direct and indirect reciprocity (with both image scoring and the standing strategy).

The standing strategy as described by Milinski *et al.* [3] is that players should not just aim for a good fitness, but also good standing. Everyone starts on a good standing, but if a player is a donor to a good standing recipient and does not cooperate with them the donor loses their good standing.

The idea of this is that committing to bad actions against good individual is immoral, but it is not necessarily true that committing to good actions towards bad individual is moral and according to most game-theoretic models it does not produce stable cooperative society. This is seen as a benefit over image scoring as with image scoring it is not always in a players interest to punish those with a low image score, as they lose reputation themselves and

in turn reduces the likelihood of them being the recipient of a cooperation. Whereas reputation is only lost in the standing strategy if the recipient is not of a good standing.

iii. Mixed Reciprocity Models

Gilbert Roberts puts forward mixed direct and indirect reciprocity models [7] using image scoring in one model similar to Nowak and Sigmund [5] and standing strategy. Both models use the island and reproductive systems Leimar and Hammerstein [2] utilized. Roberts put forward this notion due to his perception that indirect reciprocity alone is not a generalisable concept due to the close-knit nature of many societies.

Decisions can be made by agents in Roberts' model on the basis of either a reputation or an experience score. Reputation scores use image scoring in one model, differing from Nowak and Sigmund's as the score is within the range -1 to 1, initially at 0, decreasing by one for a defection and increasing by 1 for a cooperation. The standing strategy version assigns a -1 to a player when they defect against an individual or 1 to a player for anything else.

Experience scores are analogous to direct reciprocity for Roberts. When using image scoring, the experience scores between 2 players are initially 0 and go to -1 if they have defected or 1 if they have cooperated in the previous interaction. For standing strategy -1 is assigned only if the players has defected with a partner that has an experience score of 0 or 1 with them.

As you can see experience scores only reflect the previous interaction between the two individuals. This is a clear limitation as the model is seeking to describe a society made up of agents that have a higher intelligence.

Interestingly Roberts also considers the case where an agent has no resources to cooperate and thus must defect.

Steve Phelps [6] also recognises the same issue with indirect reciprocity as Gilbert Roberts - namely that remeeting in groups is likely, especially in smaller groups according to Phelps. Due to this Phelps suggests a mixed frame-




Figure 3: The spread of information through a population using indirect reciprocity with gossip and onlookers

work, to test how group sizes affect the use of direct or indirect reciprocity by agents.

iv. Gossip and Onlookers

It is recognised by Nowak and Sigmund [5] that it is unrealistic to expect all players to have directly observed a given interaction. Yet their answer - onlookers - hampers the evolution of cooperation, which relies on the probability that a player knows the image score of another player.

Sommerfeld *et al.* highlighted the importance of communication to the management of reputations. Gossip is a type of communication that is key part of societal mechanics, the structure of which is highlighted in figure 3. The findings of their experiment were that gossip needs to accurately reflect the interaction and the recipient of the gossip needs to correctly acted upon for gossip to be an adequate replacement for direct observation.

v. Comparison of Models

The aim of Roberts [7] is to compare the effectiveness of indirect reciprocity models and direct reciprocity. Roberts' results support Nowak and Sigmund's 1998 result that image scoring can support cooperation is robust even when criticisms relating to genetic drift and errors are taken into consideration.

However, it is highlighted that image scoring does not take into account who the donor is defecting against, whereas standing strategy does. Due to this image scoring suffers from an inability to distinguish between pure defectors and those who would cooperate with a cooperator. Standing strategy gives information both about how agents have interacted and the context of this action, an advantage over image scoring.

Counter to standing strategies perceived benefits it is seen that players are more susceptible to errors in perception, as standing falls faster than image score. Roberts also criticizes the standing strategy as not being a representation of the evolution of cooperation, as good standing can't have been established before the game had begun.

Roberts also points out that cooperation is only worth it when that cooperation is improving your chance of being cooperated with. This could be a counterpoint to the effectiveness of gossip as an alternative to direct observation. In gossip, there is a possibility of interactions being distorted by incorrect gossip. It could be interesting to see how the spread of misinformation could permeate society.

III. A CONCRETE MODEL

i. Specification of the Model

Both image scoring and standing strategy are good models for indirect reciprocity. They allow an agent to interpret a reputation and act as they see fit. However, the reputation of each player being universally viewable is unrealistic. The idea of onlookers seems closer to real life but is restrictive to the evolution of cooperation. The addition of gossip from these onlookers could be a sufficient alternative to direct observation, but gossip can be distorted by players wishing to influence proceedings against cooperation.

Gossip complicates a model because there is a need to interpret the information gossiped by another player. This also highlights an issue with the models that use purely standing

strategy or image scoring: reputation is not a centralised system with everyone interpreting events the same way. Reputation has a very subjective nature, especially when gossip is introduced. Due to the subjectivity of perception about an agents reputation, I do not believe that one specific model for all agents can be considered appropriate. How a player views another players reputation must be up to their inside workings.

This requires an agent to have a view of past events and also a way to interpret these events in order to make decisions about another player. A model that uses this decentralised judgement on past events easily becomes a hybrid indirect and direct reciprocity game, as agents can have their own way to interpret interactions.

For an agent to interpret gossip, the gossip must be comprehensible by them and thus should have a given structure. In our model, gossip shall be split into two categories: positive and negative. Gossip will also involve 3 agents, the agent gossiping, the recipient of the gossip, and the gossip that the agent is about (who shall not be able to view that the gossip has taken place). Gossip is intended to change how the recipient of the gossip views the agent the gossip is about - for example, if agent 2 trusts gossip from agent 3 that is negative about agent 1 and agent 2 is using the standing strategy, then agent 2 will revise their beliefs about agent 1 to have a bad standing.

The structure of the game will be as follows. There is one singular population (this is to not overly complicate the model, though changing this to an island population model such as in Leimar and Hammerstein [2]), a population is replaced by their offspring after the conclusion of a generation. A generation is made up of a number of interactions, all at a distinct time point (the number of interactions per generation is a variable).

In each time point, each agent will perceive their environment and then decide what action to take, this action may be as a donor in an interaction or maybe a piece of gossip. Each interaction shall be made up of a donor-recipient

pair, chosen at random from the population. For each interaction, a subset of the population should be chosen as direct observers (being onlookers and the donor-recipient pair).

Finally, reproduction occurs at the end of every generation and will be relative to each players fitness. The fitness of a player represents how successful they are, fitness is altered in each interaction if the donor cooperates they lose a fitness point but the recipient gains two, if the donor defects they both gain nothing and lose nothing see table 1. Reproduction represents the replacement of agents with newer agents in a network. The new agents developed for a system would likely be influenced by successful strategies in past generations thus the reproduction in our system represents the trend towards certain successful strategies in a multi-agent system.

Donor Action	Payoffs	
	Donor	Recipient
Cooperation	-1	2
Defection	0	0

Table 1: *The payoff for my indirect reciprocity model*

The group size remains the same over generations in order to measure the success of different strategies in certain group sizes, so in a new generation individuals are generated and assigned strategies. The chance an individual will be a certain strategy is worked out by summing the fitness of all individuals of that strategy and dividing it by the fitness of the overall group. Reproduction in this way has been chosen to represent the way that new agents will tend to be developed based on past successful strategies. The user will be able to select if mutation is active or not, if it is there will be a 1% chance that a new individuals strategy will be chosen completely at random.

ii. Specification of Agent Strategies

At each time point in the generation that an agent is a part of, the agent may perceive new information about interactions and gossip that has occurred in the game by way of an environment sending new percepts to agents. The agent will then be asked to decide on what they want to do at that same time point.

The percepts that the agent can receive is that they are a donor at this time point, that they are a recipient at this time point, gossip from another agent and observation of an interaction at the last time point.

Based on the percepts an agent will be able to revise their beliefs about other agents and the environment. For all agents, if they perceive that they are a donor in a donor-recipient pair they revise their beliefs to know that they are a donor at this time point. However, the revision of beliefs is generally strategy dependent, for example, in an agent using the standing strategy if they observe an interaction where a donor defects against an agent that you consider to have good standing the observing agent will revise their beliefs to give a bad standing to the donor.

Agents will have a theory on the best action to take based on their beliefs, if they have not yet received percepts they may not have any beliefs yet, so will have to have a default action. Based on the agents' beliefs and strategy they will have a theory encoded on how is best to act at each time point. The theory encoded is dependent on the strategy they are using, for example, if a discriminator (with $k=2$) has perceived they are a donor and that the recipient of their action has an image score of 1 the theory will say it is best to defect, as the image score is lower than their value of k . It will also be important to encode a theory on how to act if an agent is not a donor (how do they spread gossip? Do they remain idle?).

Of course, some agents will not care about percepts on other agents and will purely decide according to a plan such as always defect, or always cooperate. However, there needs to be a distinction on how an agent reacts to gossip

and how they react to an observation. I propose two classes of interpretation: trusting and distrusting. Distrusting agents will never take into account gossip, however, trusting agents will have a trust model, for example trusting agents using the standing strategy will only believe gossip and allow it to change their view of another agent's standing if the agent gossiping to them has a good standing from their point of view. Another example of this is if an agent is using a distrusting standing strategy, negative gossip will not change their beliefs on the standing of a certain agent.

Using this design, strategies that will be implemented include: trusting and distrusting 'neutral' standing discriminator (similar to Roberts' version), trusting and distrusting image scoring discriminator (with varying values of k), cooperator, defector. This is not an exhaustive list and I will be including more strategies based on 'neutral' standing and image scoring. These strategies only describe how to act as a donor based on beliefs on another agent. An agent will also need a strategy of how to act when it is not a donor. Some strategies may include: stay idle, always gossip accurate information to random players (if none then stay idle), always gossip inaccurate information to random people, always gossip negative or positive information to random players, randomly gossip positively or negatively. More advanced agents will have a better-developed strategy for when they are not a donor.

In conclusion, basic agents will need a way to take in percepts, they will then need a strategy to act if they are a donor and if they are not a donor. More advanced agents will use percepts to revise their beliefs and then use their beliefs to decide on how to act at any given time point. Even more advanced agents than these will develop beliefs or act on beliefs in more interesting ways than discriminating based on their view of whether an agent has a good standing or whether their view of an agent's image score is above a certain value. Examples of these more advanced strategies are: basing their action on both their own and the recipients standing or image score or using a

reinforcement algorithm to decide an action based on a belief.

IV. DISCUSSION AND CONCLUSION

Multi-agent systems require agents to be autonomous, they should not be working on centralised beliefs such as is used by Nowak and Sigmund [5] at first. My model not only removes those centralised beliefs and makes them part of the agent's inner workings but also makes the whole strategy they use part of the agent's inner workings.

The lack of individuality is where the work on indirect reciprocity I have seen falls short, because of the nature of indirect reciprocity being a community-driven mechanism the implementors see beliefs as community held and fail to realise that the beliefs may be community driven, but are actually individually held. This model is also very open to extension. Extensions on the population model (maybe using the island idea), the strategies (different interpretations of gossip and observations, strategies for the spread of gossip and strategies for actions as a donor) and mixing in indirect reciprocity by viewing observations where the agent is a recipient differently.

REFERENCES

- [1] Definition of "genetic drift" - english dictionary. <https://dictionary.cambridge.org/us/dictionary/english/genetic-drift>, 2018. Accessed: 27/11/2018.
- [2] Olof Leimar and Peter Hammerstein. Evolution of cooperation through indirect reciprocity. *Proceedings of The Royal Society*, 268:745–753, May 2001.
- [3] Theo C. M. Bakker Manfred Milinski, Dirk Semmann and Hans-Jürgen Krambeck. Cooperation through indirect reciprocity: image scoring or standing strategy? *Proceedings of The Royal Society*, 268:2495–2501, May 2001.
- [4] Martin A. Nowak. Five rules for the evolution of cooperation. *Science*, 314, 2006.
- [5] Martin A. Nowak and Karl Sigmund. Evolution of indirect reciprocity by image scoring. *Nature*, 393:573–577, 1998.
- [6] Steve Phelps. An empirical game-theoretic analysis of the dynamics of cooperation in small groups. *Journal of Artificial Societies and Social Simulation*, 2016.
- [7] Gilbert Roberts. Evolution of direct and indirect reciprocity. *Proceedings of The Royal Society*, 275:173–179, September 2008.

System design report: Multi-Agent System

JAMES KING

Supervisor: Kostas Stathis

October 2018

Abstract

Architectures and designs for intelligent agents and the environment they reside come in many different variations. For my multi-agent system I need to make some decisions to match a design for agents and their environment to the model I laid out in my report on indirect reciprocity. To do this I will consider different agent architectures, features and designs, considering how closely they match the model before laying out an architecture of my own. I shall also consider the properties of the environment the agents reside in before laying out a concrete execution for a game of indirect reciprocity. From these two component definitions I will formulate a method of communication between the environment and the agents service.

I. INTRODUCTION

One of the aims of my project is to produce a multi-agent system to play games of indirect reciprocity, the model for which I have defined in my second report: "Report on indirect reciprocity, strategies for agents and the development of a concrete model to implement". For this I have designed a system that includes two main components: The environment and a web service to host agent's decision making components (see figure 1).

I have already decided on a web service to host agent's minds for a number of reasons. One of the reasons being that in interactions across a network agents provide an API to work with (much like the API I will be providing). Meaning that it is good practice for me to consider how an API for an agent can be designed and deployed. Another reason is the difficulty interfacing between prolog and python and finally that a key use case for agents is as part of distributed systems, so using a web service to deploy them brings me closer to a real use case.

In my project, I will be hosting the environment in my Flask web application, but, the aim is for the environment to be used as if it is a




Figure 1: UML Package diagram to display the system components

library, as long as it is connected to an application/service that runs agent's decision making component. As such I will not be addressing how the web application hosting the environment, but the environment itself and the agents service.

II. CONTENT AND KNOWLEDGE

This section I will split up into 3 parts: one for each of the two main components - the environment and agents minds service and another for the communication between these two components.




Figure 2: A model-based reflex agent from Russell and Norvig [3].

i. Agents

There are many different definitions of what properties a system needs in order to be considered an agent. One definition containing 4 properties was proposed by Wooldridge and Jennings [6]: autonomy, reactivity, proactivity and social ability.

Autonomy being considered the ability to operate without direct human intervention and having control over the agent's own actions and internal state. Reactivity is defined as agents perceiving their environment and responding to changes in it in a timely fashion. Proactivity is given as the ability to exhibit goal-driven behaviour, rather than just reacting to its environment. Lastly, social ability is interaction between agents and/or humans via an ACL (agent communication language).

This notion of using the 4 properties to define an agent is considered a 'weak' notion. Wooldridge and Jennings suggest that stronger notions of agency include more human-like features. One such human-like feature that exists in game-theoretic models is the idea of memory or an internal state that the agent can act on. An internal state is used in Russell and Norvig's model-based reflex agent which takes percepts, modifies its internal state and uses the internal state and a model of the environment as input to condition-action rules as seen in figure 2. The model-based reflex agent architecture can be applied as an exam-

ple to a game-theoretic agent strategy known as an image scoring discriminator, as laid out by Nowak and Sigmund [2]. An image scoring discriminator's condition-action rules state: if the discriminator's internal state is greater than or equal to k (a variable set at the initialization of the agent) it is best to cooperate, else the agent will defect. The image scores are held in an agent's internal state and may change depending on the percepts the image scoring discriminator agent receives.

Some game-theoretic strategies don't use memory or internal state and just seek to provide a theory to act on. Though this does not qualify them as an agent according to Wooldridge and Jennings [6] they are still relevant strategies for game-theory.

These indirect reciprocity strategies bear a remarkable resemblance to both model-based reflex agents and how theories are encoded in deductive reasoning agents. Deductive reasoning agents use theories to encode how it is best to act under any given situation [5]. Agents who follow the image score discriminator strategy could have a theory encoded in them such as:

```

interaction(me, Recipient, time)      ^
image_score(me, Recipient, Score, time) ^
Score ≥ k → do(cooperate(Recipient))

interaction(me, Recipient, time)      ^
image_score(me, Recipient, Score, time) ^
Score < k → do(defect(Recipient))

¬interaction(me, _, time) → do(idle)

```

Part of the deductive reasoning agents method of implementing agents is the logical database that includes information on the current state of the world. In the example above this would possibly include logical data such as:

```

image_score(agent1, agent2, 2, 9).
interaction(agent1, agent2, 6).

```

This logical database is similar to the model-based reflex agent architecture's internal state and the human-like concept of belief. Beliefs as a mental category is included in the language Agent0 presented by Yoav Shoham [4] as one of his two mental categories: belief and commitment.

He defines beliefs as a fact that is thought to be true by an agent at a specific time about a specific time (an agent is constrained to not believe contradictory facts). Commitments are given as a commitment to act (restricted by the agent's capabilities) not a commitment to pursue a goal.

Agent capabilities are another concept Shoham uses and are defined as relations between the agent's mental state and their environment. An agent is only capable of committing to an action iff they believe themselves to be capable. In my example above this is shown by the agent only being capable of defecting or cooperating if they are a donor in an interaction.

In my system, I wish to incorporate similar ideas to those I have mentioned: beliefs (internal state), commitments, capabilities and strategies (deductive reasoning theories). Russell and Norvig specify that an agent can be thought of as a system that perceives its environment and acts within it [3]. They then go on to explain that an agent maps percept sequences to actions. In the system I am producing, this mapping will be similar to model-based reflex agents. The mapping will be provided by changing the agent's internal state based on the percepts an agent receives and then deciding on an action to take, based on this internal state/beliefs, by way of a deductive reasoning theory.

Percepts in my system will include an observation of an action (either cooperating or defecting), an observation of a gossip action (either positive or negative) and perceiving that they are a member of a donor-recipient pair at a given timepoint (and which role they hold). This raises the question: how do agents change their internal state based on the percepts they receive? The changing of the internal state is strategy dependent, but the overriding concept




Figure 3: Visual description of the process my agents' will follow

is for the agents to follow a game-theoretic approach.

To extend the example of the image scoring discriminator, let us say that agent1 perceives agent2 defecting against agent3. The image scoring strategy laid out by Nowak and Sigmund [2] states that when a defection is perceived, the perceiving agent will lower the donor's image score in their beliefs.

An observation on this method: percepts cause an agent's beliefs about its environment and other agents to change at specific timepoints. This method of changing the internal state is remarkably similar to an approach to reasoning about events (similar to percepts) and time and how events change 'fluents' (similar to beliefs) known as the events calculus [1].

Due to these remarkable similarities, the use of the event calculus seems an intuitive way to program an agent to react to percepts and query beliefs. There is a version of the event calculus known as the multi-valued fluent cached event calculus (kindly provided by my supervisor Professor Kostas Stathis) which will make querying beliefs efficient.

I have covered how agents will map from percepts to beliefs and then beliefs to actions. In each timepoint of an indirect reciprocity game, this mapping will be executed as part of a cycle step. The timepoints are executed in sequence, each cycle step includes in this order the processes perceive, revise (called as a consequence of perceiving) and decide. In the perceive step agents will receive percepts from the environment, passed through an API, from these percepts they will revise their beliefs. An agent will then decide on an action to take, the agent will commit to an action at the cycle step time-

point based on their beliefs.

The capability of an agent will be constrained at a given timepoint by way of an agent perceiving that they are the donor of a donor-recipient pair at that timepoint. The decision on which agents are donors and recipients in pairs falls to the environment. If an agent perceives they are a donor they can only cooperate or defect, but when they are not they cannot cooperate or defect. At any other point when they are not a donor they will be able to gossip or be idle.

In conclusion, I believe that the agent structure I have delineated matches the 4 properties Wooldridge and Jennings described [6]. Autonomy has been satisfied by the way that agents have control over their beliefs, control over the actions they commit to based on these beliefs and merely choose to act based on their environment, not human action. Reactivity is satisfied by the process of receiving percepts, formulating beliefs from them, and then being able to respond to the changes in the environment by acting based on these beliefs.

In my system agents have 5 options as to the actions they will be able to take: idle, positive gossip, negative gossip, cooperate as a donor and defect as a donor. Both the last two properties are met by the possibility of gossip actions, and for proactivity the ability to choose between cooperating and defecting. Gossip actions are communicative actions where a gossip communicates to another agent either positive or negative ideas about a third agent, displaying social ability. Gossip actions also show proactivity as agents can actively seek to inform other agents in a way that will affect the overall game to bring them closer to their goal.

An example of this is an agent whose goal it is to induce the evolution of cooperation gossiping to another agent a negative idea about an agent that it has viewed defecting, in the hope that the recipient of the gossip will then not help the defector by cooperating with them if they're a donor.

ii. Environment

So far I have discussed agents in my system, but agents are useless without an environment to observe and act in. My environment will run the indirect reciprocity model I outlined in my report on indirect reciprocity. The sequence of how the environment runs is given in figure 4 on page 8.

The environment consists of a community that contains 'generations' of players. The community simulates the indirect reciprocity model by creating and simulating a series of generations. Simulating a generation consists of running for each timepoint from start to end a cycle step. Each cycle step consists of 3 parts: perceive, decide and execute. Revision of beliefs is done as a consequence in an agent's mind of the perceive part, so is excluded from the environment's cycle step.

Each part of the cycle step is done for all agents at once, that is perceive sends all the percepts for a cycle step in one go, decide asks all agents to commit to a decision and then execute executes the action in the environment. Execution depends on the action if it is a gossip action or cooperate/defect action percepts are generated from this (including percepts for the selected onlookers for a cooperate/defect action) and for a cooperate action fitness is updated as per the payoff matrix in table 1.

Donor Action	Payoffs	
	Donor	Recipient
Cooperation	-1	2
Defection	0	0

Table 1: *The payoff used in my indirect reciprocity model*

The execution of an action raises the question of whether my environment is or is not deterministic. Russell and Norvig specify that when deciding on this we should ignore the actions of other agents when considering the certainty of a specific actions outcome. So when we

consider that a gossip action may or may not change the reputation of the agent it is about in the recipient's internal state this does not prevent the environment from being deterministic, as the internal state of an agent is controlled by that agent. Thus I adjudge that my environment is deterministic.

This is one of Russell and Norvig's 7 properties of a task environment, here I shall discuss further how the rest of their properties match my environment.

Due to the features of the indirect reciprocity model the environment is only partially observable. When I say the features of the indirect reciprocity model I am referring to actions which are only observable by a subset of the agents in a generation. An example of this is cooperate/defect actions which are only observable by the chosen onlookers and the donor-recipient pair.

Of course, this also highlights the environments property of being a multi-agent environment rather than a single agent environment. A more interesting question to answer on this is is the environment cooperative or competitive. This, of course, depends on whether cooperation has managed to evolve and if it is stable. But from a game-theoretic argument, even cooperative agents are competing to increase their fitness the most, the model simply encourages competition in a cooperative way. The key idea behind indirect reciprocity is: if I as a donor help this recipient, agents who have seen this will, later on, cooperate with me. Due to this key idea, my environment must be sequential, that is a decision at one timepoint may have an effect on any other decision at a later timepoint.

Russell and Norvig state that for an environment to be discrete it must handle time and actions in a discrete way. Time in the environment I am implementing will be discrete timepoints at which each cycle step executes. Actions and percepts are also discrete as they are limited to the percepts and actions mentioned in earlier sections.

I have noted previously that agents all perceive, revise and decide at the same time. The con-

sequence of this is that the environment does not change whilst the agents are deliberating over their action, thus making the environment static.

Finally, the rules of the environment are known to the agent. These rules include the payoff matrix for cooperate/defect actions, gossiping may or may not affect an agent depending on the recipient agents belief revision implementation and the idle action has no effect on the environment.

To sum up the properties of the environment I propose it is deterministic, partially observable, multi-agent, sequential, discrete, static and known.

iii. Communication and API Design

As I have mentioned previously the agents' minds will be hosted in a web service for the environment to interact with. The web service will be hosting agents' mind's and as such will need to have a creation process, this includes assigning a new mind with no previous internal state (except for any initial state which holds) and associating it with a strategy.

To follow the design principles of a RESTful system (REpresentational State Transfer) an individual agent is considered a resource, and multiple agents a collection. There is a specific path associated with the agent resource which includes a POST request associated with it to create a new agent. A POST request is also used to comply with RESTful design as it asks the server to create a resource. This is not a PUT request as the resource is not updated if the same request is attempted again, once an agent is created it cannot be changed, and thus the result is different (fails) if you repeat it.

In my system, an agent is assigned a player id, but this is not a universal identifier that distinguished it from all other agents. A player can be thought of similarly to a weak entity in a database, it is reliant on being part of a generation, which in turn is reliant on being part of a community. As such the 'primary key' of a player would be the id of the community,

generation and the player itself. This makes communities and generations resources as well which require universal resource indicators (endpoints). These endpoints both have POST requests which like the agents endpoint you can request the creation of a new resource too. When selecting agents to start a game with a user needs to know what strategies there are to select. As such, there is an endpoint associated with the strategy resource which a GET request can be sent to in order to get a list of strategies.

A cycle step combines the sequence of 4 possible steps: perceive, revise, decide and execute. Perceive and decide are mandatory steps which need to be initiated by the environment. A percept is the resource linked to the perceive step, the endpoint for this needs to be able to take new percept information and create the percept resource. Creation of the percept resource involves creating an event in the event calculus, this then may or may not change values of fluents (depending on the *initiates_at* and *causes_at* predicates associated with the event).

In the decide step an agent commits to an action, the resource that is created is an action. Action has a resource endpoint which you can send a GET request to (with parameters defining which player you are asking to decide on an action and the timepoint at which they are deciding) that returns the action that the player has decided on.

The final resource that I have so far decided that a user of the API should be the belief resource. Agents base their decisions on beliefs about other agents and their environment, and it would be great to be able to understand why an agent has made the decision that it has. A user can make a GET request to a specific type of belief (donor, recipient, interaction and standing) with parameters depending on which type of belief the user is asking for. The GET request made will return the belief of the agent at that time, which will hopefully give a clue as to why an agent acted as they did at a specific timepoint.

The API is designed to follow the 6 constraints

of a RESTful API: uniform interface (use of JSON throughout, resources with one endpoint, correct use of HTTP methods etc.), independent client and server, stateless communication, cacheable resource declaration and a layered system architecture (though for me the system will be one entity). Documentation can be found for the API in the repository of this project under the folder "FullUnit_1819_JamesKing/AgentsService/api_docs".

III. DISCUSSION AND CONCLUSION

In summation of my report, my agents will follow a similar architecture to model-based reflex agents. These agents will be able to make use of the event calculus to reason about how percepts change their beliefs. From this, they will be able to deductively reason as to what actions to commit to. These actions will be constrained by the idea of agent capabilities inspired by the Agent0 language.

The environment comprises of a number of generations, which for a certain number of cycles executes the steps perceive, decide and execute. Actions of cooperation come at a cost of 1 to the donor but a benefit of 2 to the recipient, while defections cost 1 none but gives none back. Cooperate/defect actions also generate percepts for onlookers and gossip actions for recipients.

REFERENCES

- [1] Robert Kowalski and Marek Sergot. A logic-based calculus of events. In *Foundations of knowledge base management*, pages 23–55. Springer, 1989.
- [2] Martin A. Nowak and Karl Sigmund. Evolution of indirect reciprocity by image scoring. *Nature*, 393:573–577, 1998.
- [3] Stuart J Russell and Peter Norvig. *Artificial intelligence: a modern approach*. Malaysia; Pearson Education Limited,, 2016.

- [4] Yoav Shoham. Agent0: A simple agent language and its interpreter. In *AAAI*, volume 91, page 704, 1991.
- [5] Professor Kostas Stathis. Lecture notes in intelligent agents and multi-agent systems, November 2018.
- [6] Michael Wooldridge and Nicholas R. Jennings. Intelligent agents: theory and practice. *The Knowledge Engineering Review*, 10(2):115–152, 1995.




Figure 4: Sequence diagram for the running of the environment

Web Application Testing Strategy

Part of my Final Year Project

James King

January 6, 2019

Figure 7.1